Microsoft

SQL Server 2016 CTP2 Columnstore HOL

Contents

Overview and setup

Estimated time to complete lab is 40-45 minutes.

Columnstore indexes are memory-optimized, column-oriented indexes used primarily in data warehousing scenarios, though as we will see, they can also be used for operational analytics. Columnstore indexes can be created as clustered and non-clustered index. Columnstore indexes organize data in columns to achieve high level of data compression to reduce both the storage cost and the IO. Though the compression will depend on the data, but typically you can expect 10x compression over uncompressed data 5x compression if the data was compressed with PAGE compression. Additionally, columnstore provides better query performance using batch-mode execution that processes a set of rows together for execution efficiency. Starting with SQL Server 2016, the Columnstore indexes can be used on memory-optimized tables.

This hands on lab will familiarize you with Columnstore indexes and help you see the positive performance impact they can have. It will also highlight some of the improvements in these indexes with SQL Server 2016 CTP2. In particular, you will learn:

1. A table with clustered Columnstore can now enforce referential integrity through foreign key constraints;
2. A table with clustered Columnstore tables can now have one or more non-clustered btree indexes for efficient execution of queries with equality and short-range predicates.;
3. A table with clustered or non-clustered columnstore indexes can run analytics queries efficiently.;
4. A table with non-clustered columnstore (NCCI) now allows DML operations. In earlier releases of SQL Server, a table with NCCI was only available for querying.
5. A Columnstore index on operational schema can provide efficient execution of analytics with minimal impact on OLTP workload.

At the end of this lab, you will have worked through some of the most common scenarios involved with Columnstore indexes and will have learned about some of the most significant improvements to Columnstore indexes in SQL Server 2016 CTP2.

Setting up your environment

1. Log into virtual machine environment using the following account information.

```
User: labuser
Password: Pass@word12
```

Note, if you have a monitor that supports a larger screen resolution than 1024 x 768, you can change the screen resolution for the lab to go as high as 1920 x 1080. By going to a higher screen resolution, it will be easier to use SQL Server Management Studio.

2. Right click on the desktop and click on **Screen resolution**.

3. Select **1366 x 786** (a good minimum screen size for using SSMS) and click **OK**.

4. Click **Keep Changes**.

5. The script file that contains the script to be used in this lab is located in the **C:\SQL Server 2016 CPT2 HOLs\LQS** folder.

6. The script file that contains the scripts to be used in this lab is located in the **C:\SQL Server 2016 CPT2 HOLs\Columnstore** folder.

7. Open **notepad.exe** to keep track of any information you want to save.

8. This lab will be using the database **AdventureWorks2016DW**. Open up SQL Server Management Studio and set that as your current database before running any code.

Explore Columnstore

In this exercise, you will explore a table with a clustered Columnstore index that has more than 11 million rows, set a foreign key, and run a query on the table. Let's get started!

Verify the Columnstore index and size of the table

1. Open **SQL Server Management Studio** and connect to the **.\CTP2** database engine instance.

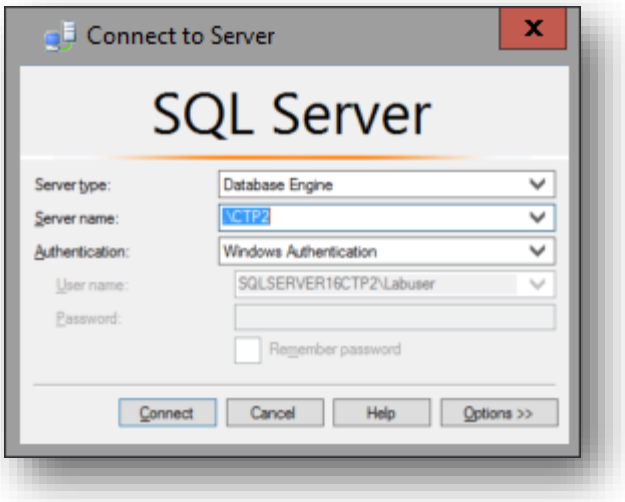

2. Expand the **Databases** node and click on **AdventureWorks2016DW**.

3. Expand the information in **Object Explorer** for the **dbo.FactResellerSalesXL_CCI** table in the **AdventureWorks2016DW** database and expand the **Indexes** folder.

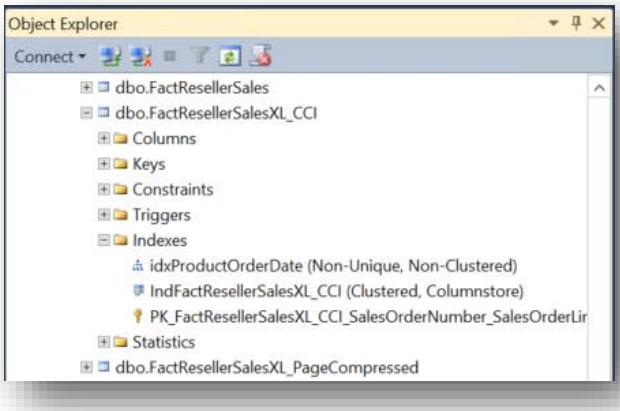

Verify that there is a CCI (clustered columnstore index) on the table.

4. Open the **C:\SQL Server 2016 CPT2 HOLs\Columnstore\Columnstore.sql** in SSMS and run the selected code below.

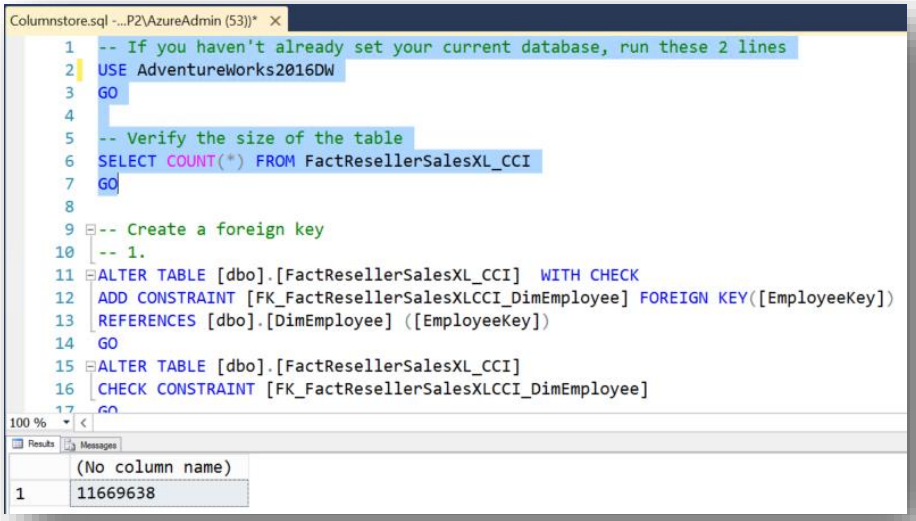

Verify that the table has more than 11 million rows.

Create a foreign key

Note that in SQL Server 2014 you could not create a foreign key relationship on a table with a Columnstore index on it.

1. Run the following two statements to create a foreign key relationship with the **DimEmployee** table.

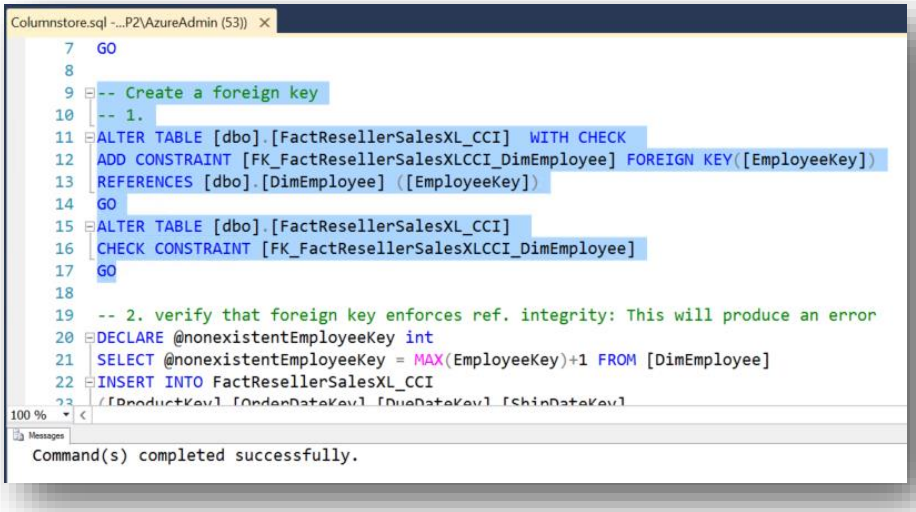

2. Verify that the foreign key constraint maintains referential integrity, by not allowing a non-existing key from the

DimEmployee table to be inserted into the **FactResellerSalesXL_CCI** table, with the following script.

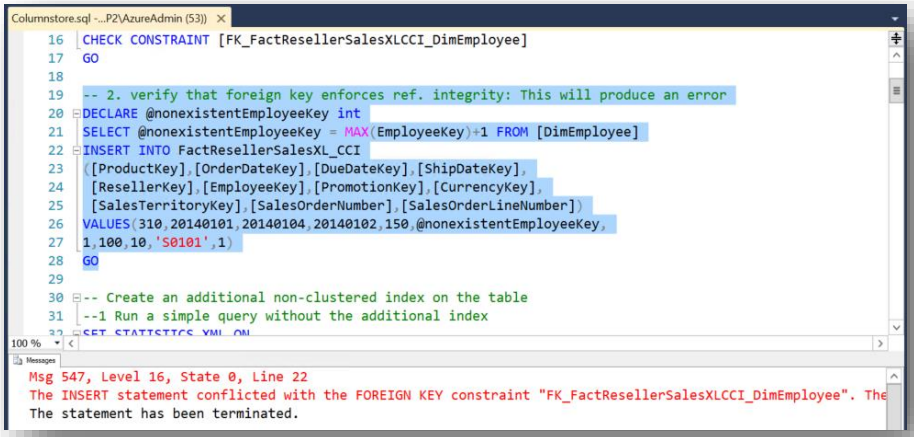

Create an additional non-clustered index on the table

Note that in SQL Server 2014 you could not create an additional index on a table that already had a clustered Columnstore index on it.

3. Run the selected query without the additional index on the table.

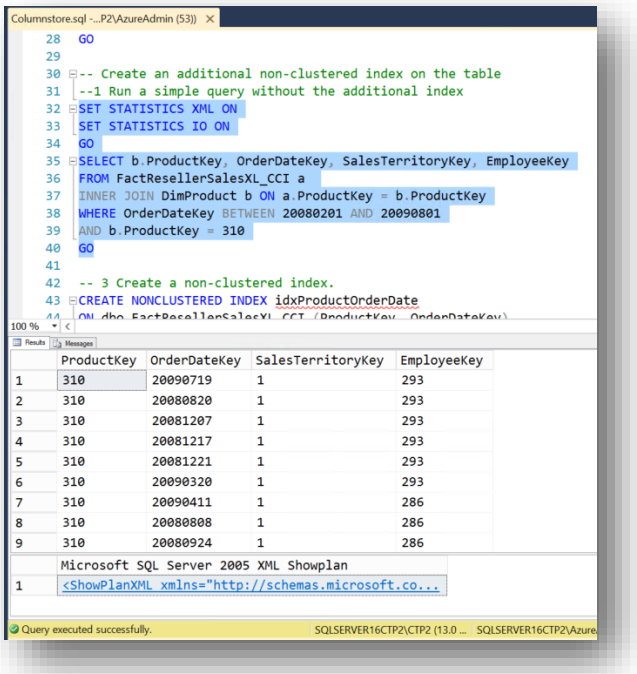

4. Click on **ShowPlanXML** link to view the query execution plan and notice the **Missing Index** hint. Also note the scanning of the entire index.

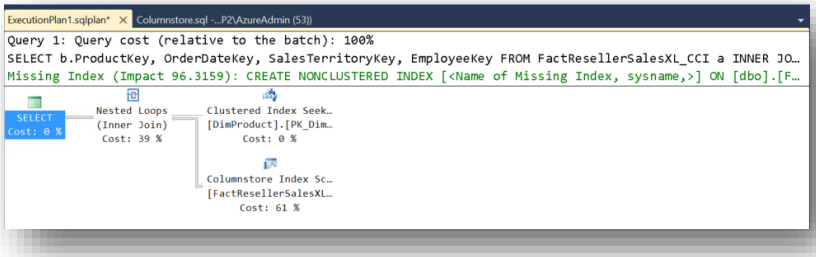

5. Go back to the **Columnstore.sql** query editor window and execute the following statement to add a new, nonclustered index. This will take a few minutes to run.

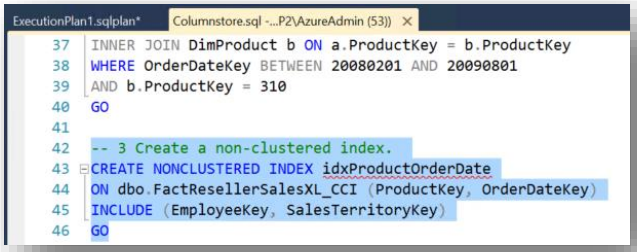

6. Now, run the above selected query again and look at the execution plan.

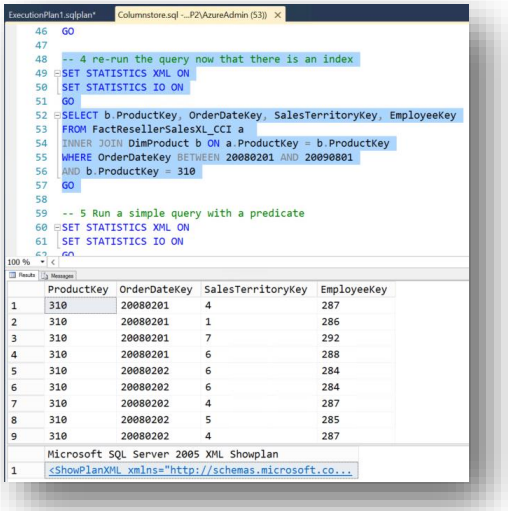

7. The index scan has now changed to an index seek.

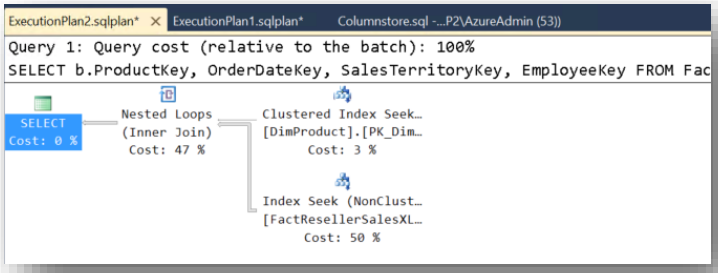

8. Let's run selected query and look at how many segments were used and how many did not have to be checked because of the index.

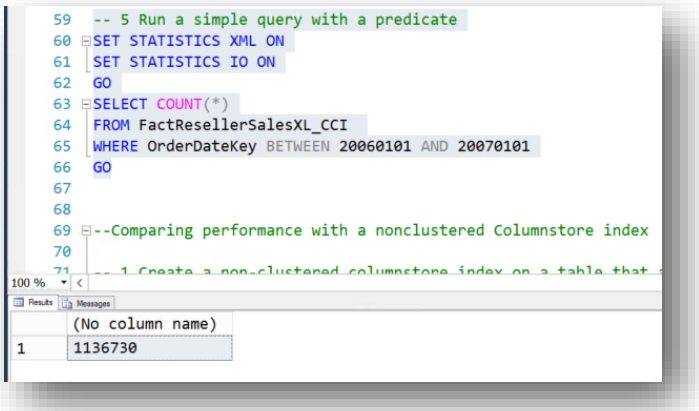

9. Look at the **Messages** tab. Notice that 12 segments were read and 2 did not have to be read. This reduces the amount of I/O that occurs, improving performance.

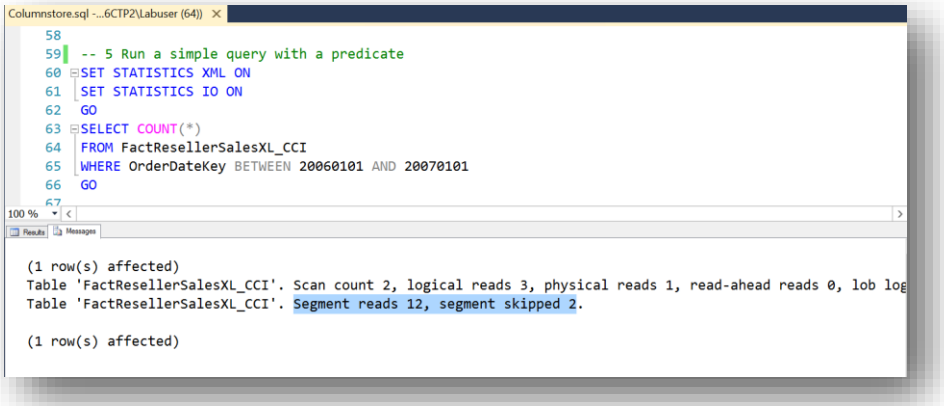

Add a Columnstore index to an existing table

In this exercise we will start with a large table, with more than 11 million rows, that is currently being used operationally. We will add a nonclustered Columnstore index to it, in order to speed up analytical queries on it. We will observe the performance impact of querying with and without the index. Then we will demonstrate that in SQL Server 2016 CTP2, tables with a nonclustered Columnstore index can now have new data inserted into them.

Comparing performance with a nonclustered Columnstore index

1. Create a new nonclustered Columnstore index on the **FactResellerSalesXL_PageCompressed** table.

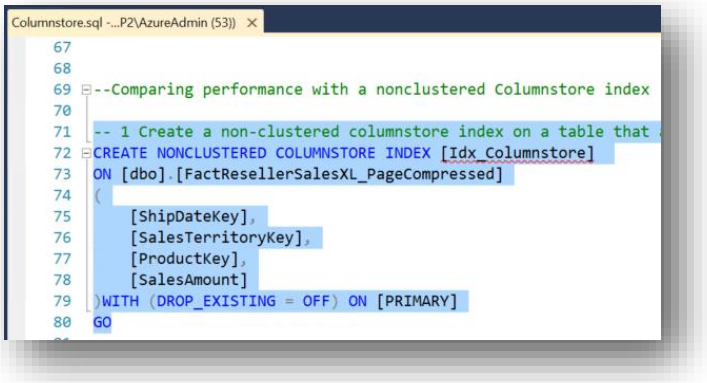

2. We can first see what the performance would have been like by adding an option to ignore the index.

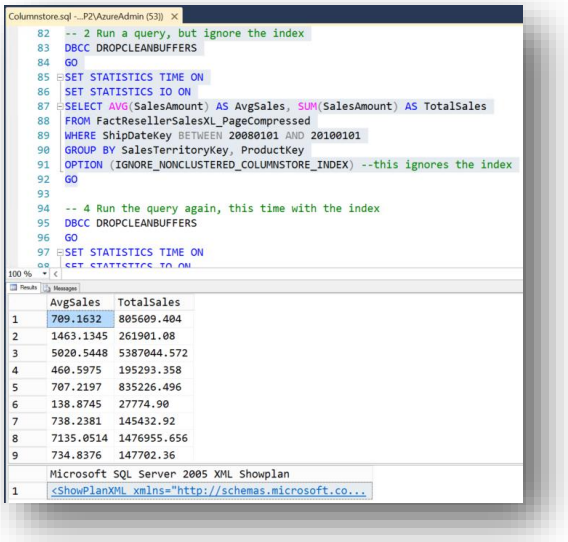

3. Run the query. After the results are returned, scroll to the bottom of the Messages tab.

```
91  OPTION (IGNORE_NONCLUSTERED_COLUMNSTORE_INDEX) --this ignores the index
92  GO
93
```

```
(3950 row(s) affected)
Table 'Worktable'. Scan count 0, logical reads 0, physical reads 0, read-ahead reads 0, lob logical reads 0,
Table 'Workfile'. Scan count 0, logical reads 0, physical reads 0, read-ahead reads 0, lob logical reads 0,
Table 'FactResellerSalesXL_PageCompressed'. Scan count 5, logical reads 87752, physical reads 1, read-ahead

(1 row(s) affected)

 SQL Server Execution Times:
   CPU time = 4703 ms,  elapsed time = 3991 ms.
```

4. Now run again without the OPTION line. Look at the Messages tab to see the performance improvement, especially in the CPU time.

```
94  -- 4 Run the query again, this time with the index
95  DBCC DROPCLEANBUFFERS
96  GO
97  SET STATISTICS TIME ON
98  SET STATISTICS IO ON
99  SELECT AVG(SalesAmount) AS AvgSales, SUM(SalesAmount) AS TotalSales
100 FROM FactResellerSalesXL_PageCompressed
101 WHERE ShipDateKey BETWEEN 20080101 AND 20100101
102 GROUP BY SalesTerritoryKey, ProductKey
103 GO
```

```
 SQL Server Execution Times:
   CPU time = 0 ms,  elapsed time = 0 ms.

(3950 row(s) affected)
Table 'FactResellerSalesXL_PageCompressed'. Scan count 8, logical reads 0, physi
Table 'FactResellerSalesXL_PageCompressed'. Segment reads 13, segment skipped 0.
Table 'Worktable'. Scan count 0, logical reads 0, physical reads 0, read-ahead r

(1 row(s) affected)

 SQL Server Execution Times:
   CPU time = 1188 ms,  elapsed time = 1291 ms.
```

5. Now look at the query execution plan. We can see that the Columnstore index is utilized and how it impacts query execution.

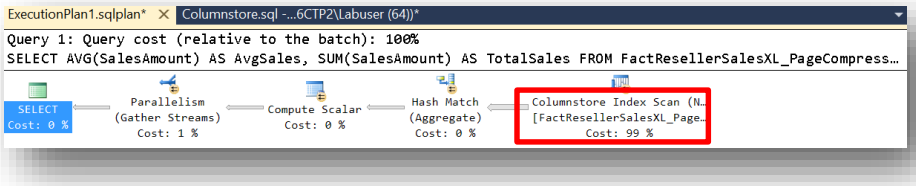

Inserting data to a table with a non-clustered Columnstore index

Note that in SQL Server 2014, it was not possible to add data to a table that had a nonclustered Columnstore index.

1. Insert 2,000 new rows by selecting the rows shown for the query below and executing the query. Verify that there is no error.

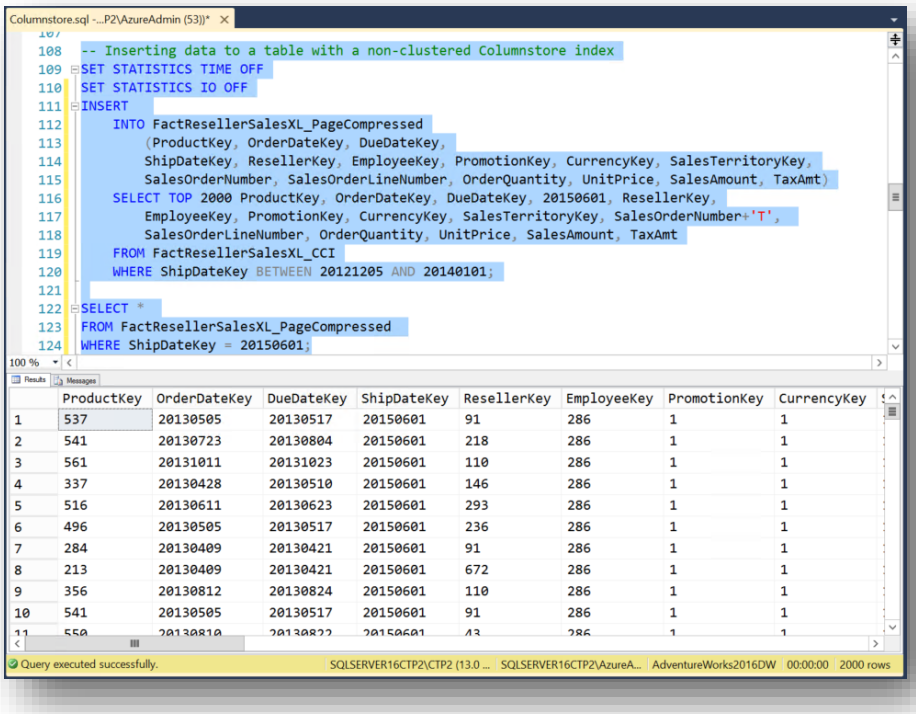

Memory-optimized tables and operational analytics

Investigate memory optimized tables

Memory-optimized tables ("Hekaton") was introduced in SQL Server 2014. We will now look at how these tables interact with Columnstore indexes and how these can be used in operational analytics.

1. Expand the **Indexes** folder in SSMS Object Explorer for the **dbo.SalesOrder** table. Notice that there are 3 indexes, including:
    - Primary key nonclustered hash,
    - Clustered Columnstore, and
    - Nonclustered indexing

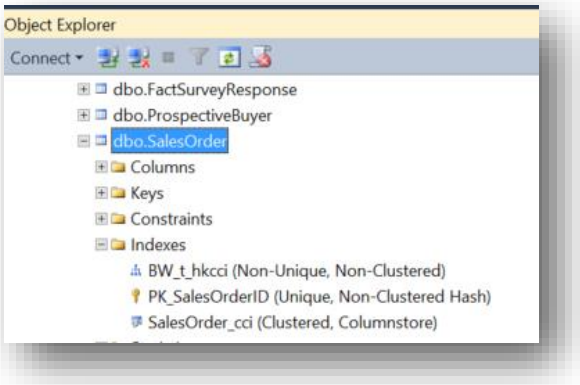

2. Also right click on the table to view its **Properties**. Notice that on the **General** page, the table's **Memory optimized** property is set to **True**.

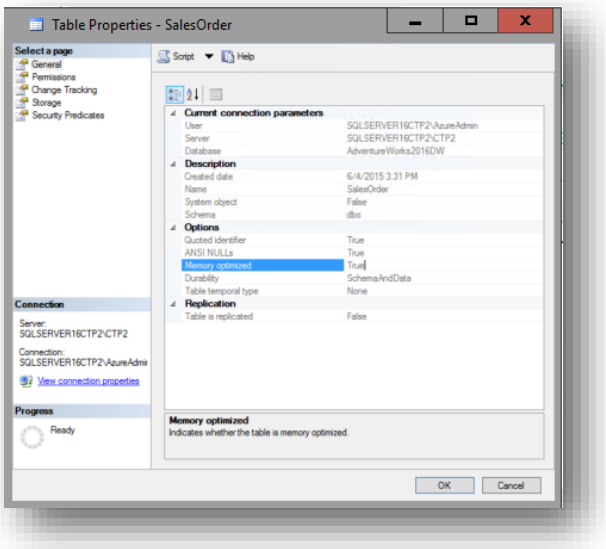

3. If you want to see the table creation script, script the table as “CREATE To.”

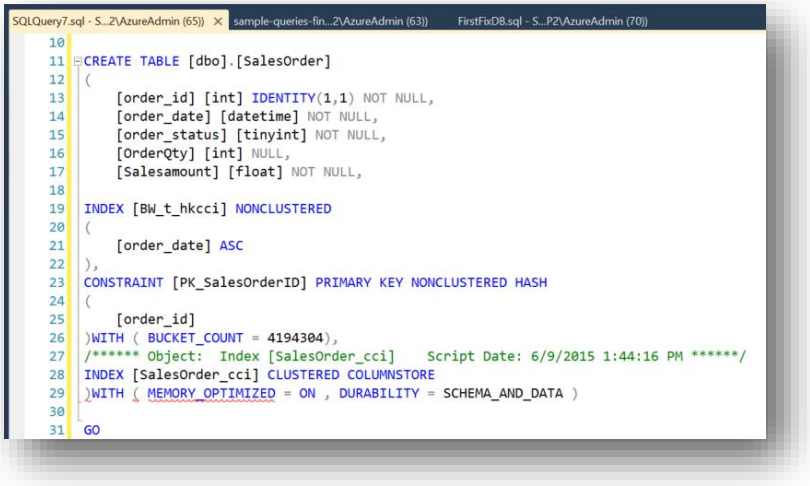

4. Now, look at the row groups information. This shows how the data is compressed.

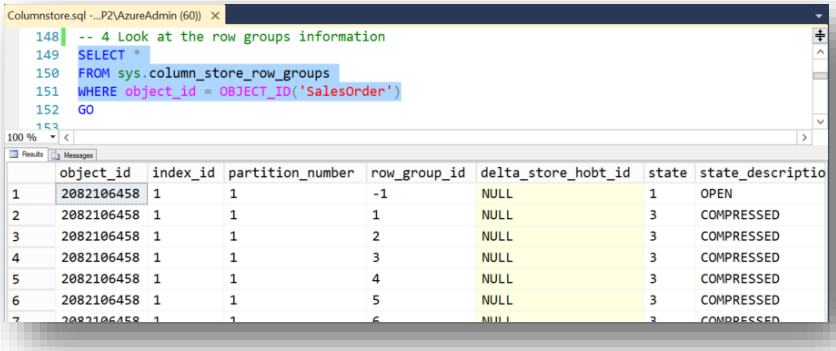

	object_id	index_id	partition_number	row_group_id	delta_store_hobt_id	state	state_descriptio
1	2082106458	1	1	-1	NULL	1	OPEN
2	2082106458	1	1	1	NULL	3	COMPRESSED
3	2082106458	1	1	2	NULL	3	COMPRESSED
4	2082106458	1	1	3	NULL	3	COMPRESSED
5	2082106458	1	1	4	NULL	3	COMPRESSED
6	2082106458	1	1	5	NULL	3	COMPRESSED

Compare performance on memory-optimized tables - count

1. Now, we’ll compare performance using the various indexes; first with a simple count query.

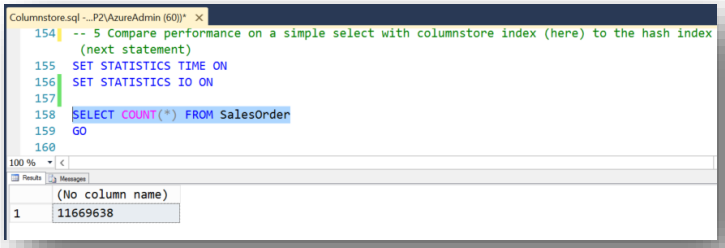

2. Look at the query execution plan. Notice that this query used the clustered Columnstore index. Also look at the time stats on the Messages tab.

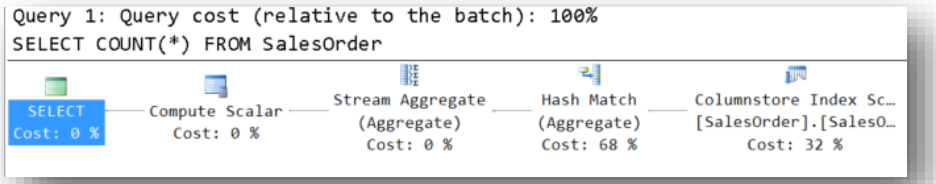

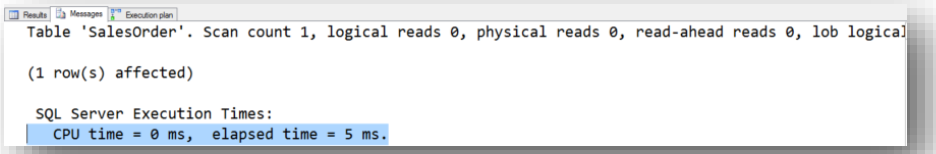

3. Now we can run the query again, but force it to use the nonclustered primary key index.

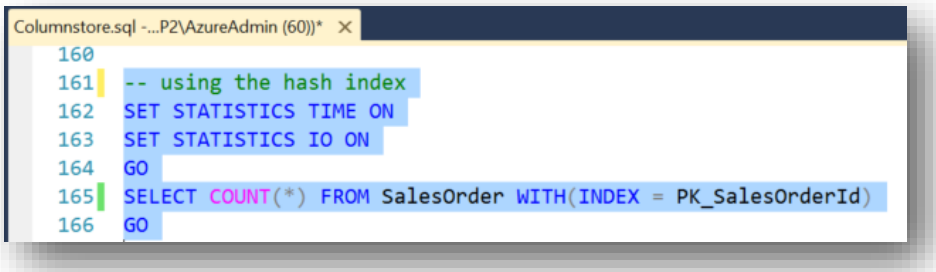

4. Again, look at the query execution plan and the time and you can see that performance is worse.

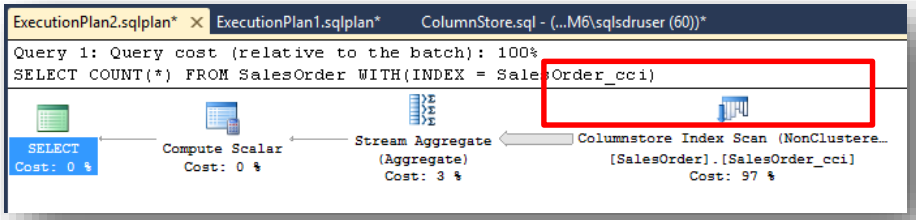

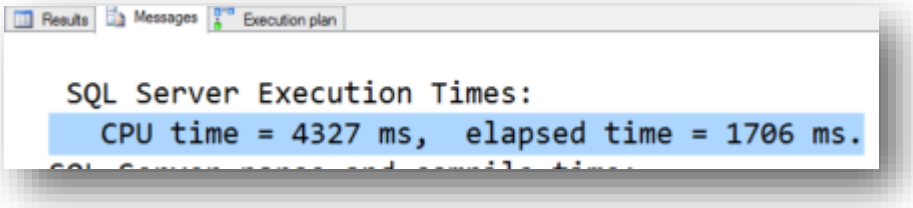

Compare performance on memory-optimized tables - analytics

The performance difference can also be significant when using analytic queries with clustered Columnstore indexes.

1. First run the query normally, so it will use its clustered columnstore index and note its time.

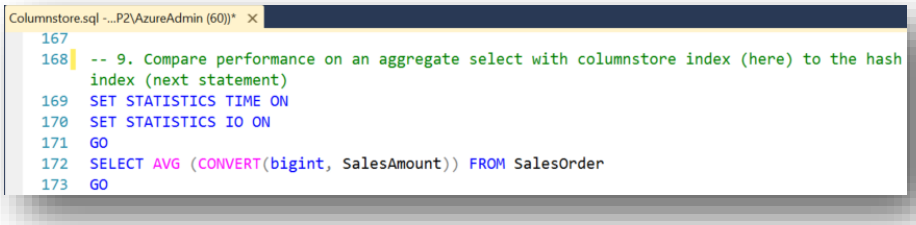

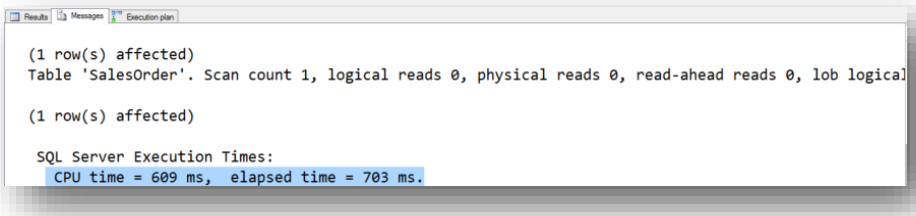

2. Now run the same query, but use the clustered Columnstore index and note the time.

```
Columnstore.sql -...P2\AzureAdmin (60))*
173  GO
174
175  -- 10. using the hash index
176  SET STATISTICS TIME ON
177  SET STATISTICS IO ON
178  GO
179  SELECT AVG (CONVERT(bigint, SalesAmount)) FROM SalesOrder  WITH(INDEX = PK_SalesOrderID)
180  GO
```

Results | Messages | Execution plan

```
SQL Server Execution Times:
   CPU time = 6251 ms,  elapsed time = 2126 ms.
```

3. We can also compare the use of a clustered Columnstore index to the performance of a natively compiled stored procedure. First the stored procedure.

```
Columnstore.sql -...P2\AzureAdmin (60))*
182  -- 11 Compare performance with a natively compiled stored proc (here) to the hash index
     (next statement)
183  SET STATISTICS TIME ON
184  SET STATISTICS IO ON
185  GO
186  IF OBJECT_ID('Hk_GetAvgSalesAmt') IS NOT NULL
187  DROP PROCEDURE Hk_GetAvgSalesAmt
188  GO
189  CREATE PROCEDURE Hk_GetAvgSalesAmt WITH SCHEMABINDING, NATIVE_COMPILATION, EXECUTE AS OWNER

190  AS
191  BEGIN ATOMIC WITH(TRANSACTION ISOLATION LEVEL = SNAPSHOT, LANGUAGE = N'ENGLISH')
192      SELECT AVG (CONVERT(bigint, SalesAmount)) FROM dbo.SalesOrder
193  END
194  GO
195  -- Execute
196  EXEC Hk_GetAvgSalesAmt
197  GO
```

Results | Messages

```
SQL Server Execution Times:
   CPU time = 2359 ms,  elapsed time = 2364 ms.
```

4. Compare this to the earlier hash index query, above.

```
Columnstore.sql -...P2\AzureAdmin (60))*
173  GO
174
175  -- 10. using the hash index
176  SET STATISTICS TIME ON
177  SET STATISTICS IO ON
178  GO
179  SELECT AVG (CONVERT(bigint, SalesAmount)) FROM SalesOrder  WITH(INDEX = PK_SalesOrderID)
180  GO
```

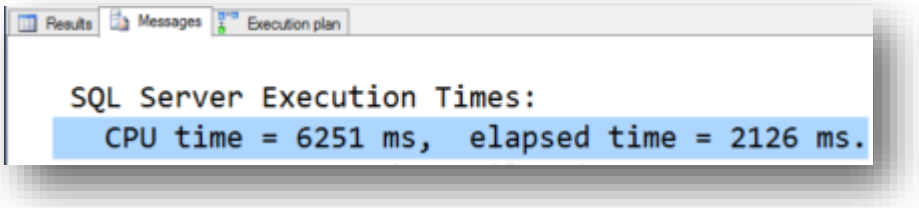

Wrap up

You should now be familiar with Columnstore indexes to see how they have a positive performance impact as described below:

1. A table with clustered Columnstore can now enforce referential integrity through foreign key constraints;
2. A table with clustered Columnstore tables can now have one or more non-clustered btree indexes for efficient execution of queries with equality and short-range predicates;
3. A table with clustered or non-clustered columnstore indexes can run analytics queries efficiently;
4. A table with non-clustered columnstore (NCCI) now allows DML operations. In earlier releases of SQL Server, a table with NCCI was only available for querying;
5. A Columnstore index on operational schema can provide efficient execution of analytics with minimal impact on OLTP workload.

Terms of use

By using this Hands-on Lab, you agree to the following terms:

The technology/functionality described in this Hands-on Lab is provided by Microsoft Corporation in a "sandbox" testing environment for purposes of obtaining your feedback and to provide you with a learning experience. You may only use the Hands-on Lab to evaluate such technology features and functionality and provide feedback to Microsoft. You may not use it for any other purpose. You may not modify, copy, distribute, transmit, display, perform, reproduce, publish, license, create derivative works from, transfer, or sell this Hands-on Lab or any portion thereof.

COPYING OR REPRODUCTION OF THE HANDS-ON LAB (OR ANY PORTION OF IT) TO ANY OTHER SERVER OR LOCATION FOR FURTHER REPRODUCTION OR REDISTRIBUTION IS EXPRESSLY PROHIBITED.

THIS HANDS-ONLAB PROVIDES CERTAIN SOFTWARE TECHNOLOGY/PRODUCT FEATURES AND FUNCTIONALITY, INCLUDING POTENTIAL NEW FEATURES AND CONCEPTS, IN A SIMULATED ENVIRONMENT WITHOUT COMPLEX SET-UP OR INSTALLATION FOR THE PURPOSE DESCRIBED ABOVE. THE TECHNOLOGY/CONCEPTS REPRESENTED IN THIS HANDS-ON LAB MAY NOT REPRESENT FULL FEATURE FUNCTIONALITY AND MAY NOT WORK THE WAY A FINAL VERSION MAY WORK. WE ALSO MAY NOT RELEASE A FINAL VERSION OF SUCH FEATURES OR CONCEPTS. YOUR EXPERIENCE WITH USING SUCH FEATURES AND FUNCITONALITY IN A PHYSICAL ENVIRONMENT MAY ALSO BE DIFFERENT.

FEEDBACK. If you give feedback about the technology features, functionality and/or concepts described in this Hands-on Lab to Microsoft, you give to Microsoft, without charge, the right to use, share and commercialize your feedback in any way and for any purpose. You also give to third parties, without charge, any patent rights needed for their products, technologies and services to use or interface with any specific parts of a Microsoft software or service that includes the feedback. You will not give feedback that is subject to a license that requires Microsoft to license its software or documentation to third parties because we include your feedback in them. These rights survive this agreement.

MICROSOFT CORPORATION HEREBY DISCLAIMS ALL WARRANTIES AND CONDITIONS WITH REGARD TO THE HANDS-ON LAB , INCLUDING ALL WARRANTIES AND CONDITIONS OF MERCHANTABILITY, WHETHER EXPRESS, IMPLIED OR STATUTORY, FITNESS FOR A PARTICULAR PURPOSE, TITLE AND NON-INFRINGEMENT. MICROSOFT DOES NOT MAKE ANY ASSURANCES OR REPRESENTATIONS WITH REGARD TO THE ACCURACY OF THE RESULTS, OUTPUT

THAT DERIVES FROM USE OF THE VIRTUAL LAB, OR SUITABILITY OF THE INFORMATION CONTAINED IN THE VIRTUAL LAB FOR ANY PURPOSE.

DISCLAIMER

This lab contains only a portion of new features and enhancements in Microsoft SQL Server 2016 CTP2. Some of the features might change in future releases of the product. In this exercise, you will learn about some, but not all, new features.